



ANCIENT METAGENOMICS

Physalia course, online, 2-4 June 2026

Read-based taxonomic profiling

Nikolay Oskolkov, Metabolic Research Group leader at NIRI and Lund University



Physalia
Courses

Typical analysis methods used in metagenomics

1) Alignment:



2) Classification:




Centrifuge

MetaPhlan

Clark

Reference based:
assume similarity to reference

3) De-novo assembly:




Reference free:
unbiased but challenging

What is competitive mapping and why you should do it

1) Single genome (mammoth) aDNA mapping:



2) Correcting for human contamination:




3) Ancient metagenomics mapping:



Competitive mapping is absolutely central for metagenomic analysis!

Evolution of taxonomic profilers (my view)



K-mer based taxonomic profiling: Kraken family of tools

Wood et al. *Genome Biology* (2019) 20:257
<https://doi.org/10.1186/s13059-019-1891-0>

Genome Biology

SHORT REPORT

Open Access

Improved metagenomic analysis with Kraken 2

Derrick E. Wood^{1,2}, Jennifer Lu^{2,3} and Ben Langmead^{1,2*}

Abstract

Although Kraken's k-mer-based approach provides a fast taxonomic classification of metagenomic sequence data, its large memory requirements can be limiting for some applications. Kraken 2 improves upon Kraken 1 by reducing memory usage by 85%, allowing greater amounts of reference genomic data to be used, while maintaining high accuracy and increasing speed fivefold. Kraken 2 also introduces a translated search mode, providing increased sensitivity in viral metagenomics analysis.


Keywords: Metagenomics, Metagenomics classification, Microbiome, Probabilistic data structures, Alignment-free methods, Minimizers

Assigning taxonomic labels to sequencing reads is an important part of many computational genomics pipelines for metagenomics projects. Recent years have seen several approaches to accomplish this task in a time-efficient manner [1–3]. One such tool, Kraken [4], uses a memory-intensive algorithm that associates short genomic substrings (k-mers) with the lowest common ancestor (LCA) taxa. Kraken and related tools like KrakenUniq [5] have proven highly efficient and accurate in independent tool comparisons [6, 7]. But Kraken's high memory requirements force many researchers to either use a reduced-sensitivity MiniKraken database [8, 9] or to build and use many indexes over subsets of the reference sequences [10, 11]. Its memory requirements can easily exceed 100 GB [7], especially when the reference data includes large eukaryotic genomes [12, 13]. Here, we introduce Kraken 2, which provides a major reduction in memory usage as well as faster classification, a spaced seed searching scheme, a translated search mode for matching in amino acid space, and continued compatibility with the Bracken [14] species-level sequence abundance estimation algorithm.

Kraken 2 addresses the issue of large memory requirements through two changes to Kraken 1's data

structures and algorithms. While Kraken 1 used a sorted list of k-mer/LCA pairs indexed by minimizers [15], Kraken 2 introduces a probabilistic, compact hash table to map minimizers to LCAs. This table uses one third of the memory of a standard hash table, at the cost of some specificity and accuracy. Additionally, Kraken 2 only stores minimizers (of length ℓ , $\ell \leq k$) from the reference sequence library in its data structure, whereas Kraken 1 stored all k-mers. This change means that, during classification, the minimizer (ℓ -mer) is the substring compared against a reference set in Kraken 2, while Kraken 1 compared k-mers (Fig. 1a, b). Kraken 2's index for a specific reference database with 9.1 Gbp of genomic sequences uses 10.6 GB of memory when classifying. Kraken 1's index for the same reference uses 72.4 GB of memory for classification (Fig. 2a, Additional file 1: Table S1). In general, a Kraken 2 database is about 85% smaller than a Kraken 1 database over the same references (Additional file 2: Figure S1).

Kraken 2's approach is faster than Kraken 1's because only distinct minimizers from the query (read) trigger accesses to the hash table. A similar minimizer-based approach has proven useful in accelerating read alignment [16]. Kraken 2 additionally provides a hash-based subsampling approach that reduces the set of minimizer/LCA pairs included in the table, allowing the user to specify a target hash table size; smaller hash tables yield lower memory usage and higher classification



* Correspondence: langmead@cs.jhu.edu

¹Department of Computer Science, Whiting School of Engineering, Johns Hopkins University, Baltimore, MD, USA

²Center for Computational Biology, Johns Hopkins University, Baltimore, MD, USA

Full list of author information is available at the end of the article



Physalia
Courses

© The Author(s). 2019 **Open Access** This article is distributed under the terms of the Creative Commons Attribution 4.0 International License (<http://creativecommons.org/licenses/by/4.0/>), which permits unrestricted use, distribution, and reproduction in any medium, provided you give appropriate credit to the original author(s) and the source, provide a link to the Creative Commons license, and indicate if changes were made. The Creative Commons Public Domain Dedication waiver (<http://creativecommons.org/publicdomain/zero/1.0/>) applies to the data made available in this article, unless otherwise stated.

Advantage of classification over alignment: speed, Kraken2 is very fast!

Why exactly do we need LCA?

Same genus

```
sequence 1 ATGGTCGGGCAGGACGTTGCGAGT
sequence 2 CGAGAAAGGGCAGGACGCCACGTAC
```

Species 1


Species 2

Ambiguous reads

Ignore ambiguous reads
(lose many reads)

Align genus reads
to your species of
interest ref genome
(isn't it too risky?)

Keep them for quantifying
abundance on genus level
(LCA)



Activities

LibreOffice Calc

File Edit View Insert Format Styles Sheet Data Tools Window Help

Liberation Sans 10

% 0.0

J34

	A	B	C	D	E	F	G	H	I	J
1	% of reads	reads	taxReads	n_minimizer	n_distinct_minimizer	rank	taxID	Name		
2	85.76	229988348	229988348	0		0 U	0	Unclassified		
3	14.24	38188511	317329	229258315	10956342 R		1	root		
4	14.1	37814928	553677	223347668	10689204 R1		131567	cellular organisms		
5	13.35	35791835	861918	201695602	8612373 D		2	Bacteria		
6	10.84	29060274	1196943	151278879	5085202 P		1224	Proteobacteria		
7	8.22	22056723	407757	106202766	1185104 C		28216	Betaproteobacteria		
8	7.82	20984490	1268537	97499372	1030255 O		80840	Burkholderiales		
9	6	16081257	2224086	68602917	538180 F		80864	Comamonadaceae		
10	2.23	5974652	286945	22944027	90448 G		52972	Polaromonas		
11	1.11	2971020	18648	10967489	38279 S		216465	Polaromonas naphthalenivorans		
12	1.1	2952372	2952372	10917548	38245 S1		365044	Polaromonas naphthalenivorans CJ2		
13	0.95	2560665	1058948	9078851	43621 G1		2638319	Unclassified Polaromonas		
14	0.55	1465800	1465800	4741406	23365 S		296591	Polaromonas sp. JS666		
15	0	8327	8327	23317	17 S		1840289	Polaromonas sp. H4N		
16	0	7639	7639	40303	221 S		1840293	Polaromonas sp. H6N		
17	0	7127	7127	35955	154 S		1840301	Polaromonas sp. W10N		
18	0	3429	3429	14648	74 S		1840297	Polaromonas sp. H8N		
19	0	2336	2336	10159	219 S		1840303	Polaromonas sp. W11N		
20	0	1757	1757	7805	272 S		1869339	Polaromonas sp.		
21	0	1251	1251	2288	2 S		1840287	Polaromonas sp. H3N		
22	0	956	956	4451	96 S		1840283	Polaromonas sp. H1N		
23	0	770	770	1317	8 S		416605	Polaromonas sp. A10		
24	0	721	721	1823	1 S		1840281	Polaromonas sp. H12N		
25	0	628	628	712	5 S		2268087	Polaromonas sp. SP1		
26	0	255	255	812	64 S		1840267	Polaromonas sp. E5S		
27	0	245	245	745	49 S		1840265	Polaromonas sp. E3S		
28	0	222	222	1167	24 S		1840239	Polaromonas sp. E10S		
29	0	163	163	405	2 S		480424	Polaromonas sp. GM1		
30	0	20	20	374	1 S		1840257	Polaromonas sp. E19S		
31	0	15	15	51	5 S		642193	Polaromonas sp. RB76		
32	0	15	15	57	20 S		1840275	Polaromonas sp. E9S		
33	0	9	9	28	16 S		1840323	Polaromonas sp. W9N		
34	0	8	8	8	1 S		1705699	Polaromonas sp. 277		

Krestovka.merged.trimmed.fastq.gz_kraken2

Find

Find All

Formatted Display

Match Case

Sheet 1 of 1

Default


English (USA)

Average: ; Sum: 0


140%

Ways to visualize and interpret taxonomic profilers outputs

KRONA



PAVIAN



How you see false-positives: Kraken2 with NT database

	A	B	C	D	E	F
1	Percent Reads	Reads	TaxReads	Rank	TaxID	Name
2	0.12	6438	6438	S	1009846	Burkholderia cepacia GG4
3	0.09	4683	4683	S	1395570	Burkholderia cepacia JBK9
4	0.09	4685	4685	S	1417228	Paraburkholderia phytofirmans OLGA172
5	0.05	2483	2483	S	1249668	Burkholderia ubonensis MSMB22
6	0.05	2736	2736	S	339670	Burkholderia ambifaria AMMD
7	0.04	2193	2193	S	391038	Paraburkholderia phymatum STM815
8	0.04	2317	2317	S	9615	Canis lupus familiaris
9	0.03	1738	1738	S	754502	Paraburkholderia sprentiae WSM5005
10	0.03	1640	1640	S	1229205	Paraburkholderia phenoliruptrix BR3459a
11	0.03	1664	1664	S	1416914	Pandoraea pnomenusa 3kgm
12	0.03	1785	1785	S	1112204	Gordonia polyisoprenivorans VH2
13	0.02	871	871	S	395019	Burkholderia multivorans ATCC 17616
14	0.02	919	919	S	398577	Burkholderia ambifaria MC40-6
15	0.02	955	955	S	398527	Paraburkholderia phytofirmans PsJN
16	0.02	982	982	S	266265	Paraburkholderia xenovorans LB400
17	0.02	973	973	S	272630	Methylobacterium extorquens AM1
18	0.02	1143	1143	S	223781	Aquila chrysaetos chrysaetos
19	0.02	1192	1192	S	202946	Apteryx mantelli mantelli
20	0.01	755	755	S	406425	Burkholderia cenocepacia MC0-3
21	0.01	623	623	S	32009	Burkholderia gladioli pv. gladioli
22	0.01	581	581	S	999541	Burkholderia gladioli BSR3
23	0.01	772	772	S	1323664	Paraburkholderia caribensis MBA4
24	0.01	536	536	S	264198	Cupriavidus pinatubonensis JMP134
25	0.01	752	752	S	882378	Paraburkholderia rhizoxinica HKI 454
26	0.01	546	546	S	1034807	Flavobacterium branchiophilum FL-15
27	0.01	587	587	S	112262	Ovis canadensis canadensis

Look like reasonable microbes so far ...



This was a blank!

Give me one best metagenomic taxonomic classifier



Selection of Appropriate Metagenome Taxonomic Classifiers for Ancient Microbiome Research

Irina M. Velsko,^{a*} Laurent A. F. Frantz,^{a,b} Alexander Herbig,^c Greger Larson,^d Christina Warinner^{c,d,e}

^aPalaeogenomics and Bio-Archaeology Research Network, Research Laboratory for Archaeology and the History of Art, University of Oxford, Oxford, United Kingdom
^bSchool of Biological and Chemical Sciences, Queen Mary University of London, London, United Kingdom
^cDepartment of Archaeogenetics, Max Planck Institute for the Science of Human History, Jena, Germany
^dDepartment of Anthropology, University of Oklahoma, Norman, Oklahoma, USA
^eDepartment of Periodontology, University of Oklahoma Health Sciences Center, Oklahoma City, Oklahoma, USA

ABSTRACT Metagenomics enables the study of complex microbial communities from myriad sources, including the remains of oral and gut microbiota preserved in archaeological dental calculus and paleofeces, respectively. While accurate taxonomic assignment is essential to this process, DNA damage characteristic of ancient samples (e.g., reduction in fragment size and cytosine deamination) may reduce the accuracy of read taxonomic assignment. Using a set of *in silico*-generated metagenomic data sets, we investigated how the addition of ancient DNA (aDNA) damage patterns influences microbial taxonomic assignment by five widely used profilers: QIIME/UCLUST, MetaPhlAn2, MIDAS, CLARK-S, and MALT. *In silico*-generated data sets were designed to mimic dental plaque, consisting of 40, 100, and 200 microbial species/strains, both with and without simulated aDNA damage patterns. Following taxonomic assignment, the profiles were evaluated for species presence/absence, relative abundance, alpha diversity, beta diversity, and specific taxonomic assignment biases. Unifrac metrics indicated that both MIDAS and MetaPhlAn2 reconstructed the most accurate community structure. QIIME/UCLUST, CLARK-S, and MALT had the highest number of inaccurate taxonomic assignments; false-positive rates were highest by CLARK-S and QIIME/UCLUST. Filtering out species present at <0.1% abundance greatly increased the accuracy of CLARK-S and MALT. All programs except CLARK-S failed to detect some species from the input file that were in their databases. The addition of ancient DNA damage resulted in minimal differences in species detection and relative abundance between simulated ancient and modern data sets for most programs. Overall, taxonomic profiling biases are program specific rather than damage dependent, and the choice of taxonomic classification program should be tailored to specific research questions.

IMPORTANCE Ancient biomolecules from oral and gut microbiome samples have been shown to be preserved in the archaeological record. Studying ancient microbiome communities using metagenomic techniques offers a unique opportunity to reconstruct the evolutionary trajectories of microbial communities through time. DNA accumulates specific damage over time, which could potentially affect taxonomic classification and our ability to accurately reconstruct community assemblages. It is therefore necessary to assess whether ancient DNA (aDNA) damage patterns affect metagenomic taxonomic profiling. Here, we assessed biases in community structure, diversity, species detection, and relative abundance estimates by five popular metagenomic taxonomic classification programs using *in silico*-generated data sets with and without aDNA damage. Damage patterns had minimal impact on the taxonomic profiles produced by each program, while false-positive rates and biases were intrinsic to each program. Therefore, the most ap-

Received 29 May 2018 Accepted 20 June 2018 Published 17 July 2018
Citation: Velsko IM, Frantz LAF, Herbig A, Larson G, Warinner C (2018) Selection of appropriate metagenomic taxonomic classifiers for ancient microbiome research. *mSystems* 3:e00080-18. <https://doi.org/10.1128/mSystems.00080-18>
Editor: Thomas J. Shapton, Oregon State University
Copyright: © 2018 Velsko et al. This is an open-access article distributed under the terms of the Creative Commons Attribution 4.0 International license.
Address correspondence to Christina Warinner, warinner@hihmpg.de.
* Present address: Irina M. Velsko, Department of Biological Sciences, Clemson University, Clemson, South Carolina, USA.
† Taxonomic classification of ancient metagenomes is minimally affected by DNA damage patterns

Submitted 12 April 2021
Accepted 21 December 2021
Published 24 March 2022

Corresponding authors
samuel.neuenschwander@unil.ch
anna-sapfo.malaspinas,
annasapfo.malaspinas@unil.ch

Academic editor
Joseph Gillespie

Additional Information and
Declarations can be found on
page 26

DOI: 10.7717/peerj.12784

© Copyright
2022 Arizmendi Cárdenas et al.

Distributed under
Creative Commons CC-BY 4.0

OPEN ACCESS

Benchmarking metagenomics classifiers on ancient viral DNA: a simulation study

Yami Omarr Arizmendi Cárdenas^{1,2}, Samuel Neuenschwander^{1,3} and Anna-Sapfo Malaspinas^{1,2}

¹ Department of Computational Biology, University of Lausanne, Lausanne, Switzerland
² Swiss Institute of Bioinformatics, Lausanne, Switzerland
³ Vital-IT, Swiss Institute of Bioinformatics, Lausanne, Switzerland

ABSTRACT

Owing to technological advances in ancient DNA, it is now possible to sequence viruses from the past to track down their origin and evolution. However, ancient DNA data is considerably more degraded and contaminated than modern data making the identification of ancient viral genomes particularly challenging. Several methods to characterise the modern microbiome (and, within this, the virome) have been developed; in particular, tools that assign sequenced reads to specific taxa in order to characterise the organisms present in a sample of interest. While these existing tools are routinely used in modern data, their performance when applied to ancient microbiome data to screen for ancient viruses remains unknown. In this work, we conducted an extensive simulation study using public viral sequences to establish which tool is the most suitable to screen ancient samples for human DNA viruses. We compared the performance of four widely used classifiers, namely Centrifuge, Kraken2, DIAMOND and MetaPhlAn2, in correctly assigning sequencing reads to the corresponding viruses. To do so, we simulated reads by adding noise typical of ancient DNA to a set of publicly available human DNA viral sequences and to the human genome. We fragmented the DNA into different lengths, added sequencing error and C to T and G to A deamination substitutions at the read termini. Then we measured the resulting sensitivity and precision for all classifiers. Across most simulations, more than 228 out of the 233 simulated viruses were recovered by Centrifuge, Kraken2 and DIAMOND, in contrast to MetaPhlAn2 which recovered only around one third. Overall, Centrifuge and Kraken2 had the best performance with the highest values of sensitivity and precision. We found that deamination damage had little impact on the performance of the classifiers, less than the sequencing error and the length of the reads. Since Centrifuge can handle short reads (in contrast to DIAMOND and Kraken2 with default settings) and since it achieves the highest sensitivity and precision at the species level across all the simulations performed, it is our recommended tool. Regardless of the tool used, our simulations indicate that, for ancient human studies, users should use strict filters to remove all reads of potential human origin. Finally, we recommend that users verify which species are present in the database used, as it might happen that default databases lack sequences for viruses of interest.

How to cite this article: Arizmendi Cárdenas YO, Neuenschwander S, Malaspinas A-S. 2022. Benchmarking metagenomics classifiers on ancient viral DNA: a simulation study. *PeerJ* 10:e12784. <https://doi.org/10.7717/peerj.12784>

Benchmarking Metagenomic Classifiers on Simulated Ancient and Modern Metagenomic Data

Vaidehi Pusadkar^{1,2} and Rajeev K. Azad^{1,2,3,*}

¹ Department of Biological Sciences, University of North Texas, Denton, TX 76203, USA;
² vaidehipusadkar@myunt.edu
³ BioDiscovery Institute, University of North Texas, Denton, TX 76203, USA
⁴ Department of Mathematics, University of North Texas, Denton, TX 76203, USA
* Correspondence: rajeev.azad@unt.edu

Abstract: Taxonomic profiling of ancient metagenomic samples is challenging due to the accumulation of specific damage patterns on DNA over time. Although a number of methods for metagenome profiling have been developed, most of them have been assessed on modern metagenomes or simulated metagenomes mimicking modern metagenomes. Further, a comparative assessment of metagenome profilers on simulated metagenomes representing a spectrum of degradation depth, from the extremity of ancient (most degraded) to current or modern (not degraded) metagenomes, has not yet been performed. To understand the strengths and weaknesses of different metagenome profilers, we performed their comprehensive evaluation on simulated metagenomes representing human dental calculus microbiome, with the level of DNA damage successively raised to mimic modern to ancient metagenomes. All classes of profilers, namely, DNA-to-DNA, DNA-to-protein, and DNA-to-marker comparison-based profilers were evaluated on metagenomes with varying levels of damage simulating deamination, fragmentation, and contamination. Our results revealed that, compared to deamination and fragmentation, human and environmental contamination of ancient DNA (with modern DNA) has the most pronounced effect on the performance of each profiler. Further, the DNA-to-DNA (e.g., Kraken2, Bracken) and DNA-to-marker (e.g., MetaPhlAn) based profiling approaches showed complementary strengths, which can be leveraged to elevate the state-of-the-art of ancient metagenome profiling.

Keywords: microorganisms; microbial DNA; ancient metagenomics; taxonomic profiling; benchmarking

1. Introduction

The evolutionary history of genetic information of all extant microorganisms goes back up to billions of years in time. Advances in metagenome sequencing have enabled the reconstruction of evolutionary trajectories to understand how microbial communities have evolved through time. Such studies have highlighted the emergence and evolution of microbes involved in infectious diseases [1,2], ancient pandemics [3], and migrations that happened in the past. Ancient metagenomics refers to the analysis of complex genomic content which is retrieved from degraded biological material of ancient or extinct microorganisms. Research in this field is commonly focused on adding a new perspective to our understanding of the past and providing exciting sources of information for modern metagenomics studies. With the decrease in sequencing cost and increase in coverage depth, thanks to the next-generation high throughput sequencing technologies, reliable estimation of the relative abundance of organisms (or taxa) dwelling in an environment is now possible [4]. Popular methods for profiling microbial communities use sequence data obtained using either 16S rRNA amplicon sequencing or whole shotgun metagenome sequencing. The former targets the 16S rRNA gene for taxonomic profiling, while the latter targets the entire complement of genomic information in a community to profile

Microorganisms 2023, 11, 2478. <https://doi.org/10.3390/microorganisms11102478>

<https://www.mdpi.com/journal/microorganisms>

Are we comparing classifiers or filtering strategies of their outputs?

Why to do filtering of Kraken output?

% reads	taxReads	kmers	dup	cov	taxID	rank	taxName
87.090	222111	3577712	1.00	9.593e-03	0	no rank	unclassified
12.910	32934	331	329116	1.17	3.026e-06	1	no rank
12.780	32603	3	324663	1.16	3.032e-06	131567	no rank
7.002	17859	7	111728	1.43	1.388e-06	2759	superkingdom
6.998	17849	0	111664	1.43	1.729e-06	33154	clade
6.998	17847	0	111664	1.43	1.918e-06	33208	kingdom
6.998	17847	0	111664	1.43	1.919e-06	6072	clade
6.998	17847	0	111664	1.43	1.935e-06	33213	clade
6.996	17844	0	111664	1.43	2.365e-06	33511	clade
6.996	17844	0	111664	1.43	2.392e-06	7711	phylum
6.996	17844	0	111664	1.43	2.399e-06	89593	subphylum
6.996	17844	0	111664	1.43	2.399e-06	7742	clade
6.996	17844	1	111664	1.43	2.402e-06	7776	clade
6.996	17843	0	111640	1.43	2.410e-06	117570	clade
6.996	17843	12	111640	1.43	2.410e-06	117571	clade
6.973	17785	0	111127	1.43	4.908e-06	8287	superclass
6.973	17785	0	111127	1.43	4.920e-06	1338369	clade
6.973	17785	3	111127	1.43	4.920e-06	32523	clade
6.972	17781	47	111095	1.43	5.185e-06	32524	clade
6.945	17714	1	110030	1.44	7.432e-06	40674	class

rank = species



% reads	taxReads	kmers	dup	cov	taxID	rank	taxName
4.7660000	12155	5002	72665	1.45	1.342e-03	9785	species
0.0019600	5	0	48	1.00	1.719e-03	99490	species
0.0337200	86	0	335	1.19	3.026e-06	1	no rank
0.0003921	1	0	8	1.00	3.032e-06	131567	no rank
0.3937000	1004	9	3984	1.08	1.020e-04	9778	species
0.0172500	44	0	130	1.30	1.020e-04	9778	species
0.1380000	352	25	1388	1.14	6.428e-04	9813	species
0.0003921	1	0	4	1.00	8.105e-07	9612	species
0.0200000	51	3	188	1.14	7.836e-07	34882	species
0.0047050	12	0	33	1.21	8.168e-07	9606	species
0.0019600	5	1	38	1.00	1.212e-06	9606	species
0.0003921	1	0	6	1.00	5.993e-03	1491	species
0.0011760	3	0	12	1.00	5.993e-03	1491	species
0.0003921	1	1	2	1.00	3.333e-04	237658	species
0.1313000	335	3	1351	1.01	8.105e-07	9612	species
0.0294100	75	2	277	1.06	5.837e-07	9657	species
0.0007842	2	0	10	1.00	7.836e-07	34882	species
0.0007842	2	1	10	1.00	8.168e-07	76717	species
0.0019600	5	0	16	1.00	8.997e-07	9715	species
0.0011760	3	1	10	1.00	4.989e-07	9720	species

kmers = 1000



taxReads = 200



% reads	taxReads	kmers	dup	cov	taxID	rank	taxName
4.7660	12155	5002	72665	1.45	1.342e-03	9785	species
0.1686	430	430	1794	1.00	1.212e-06	9606	species
5.1680	13180	12862	178949	1.01	5.993e-03	1491	species

% reads	taxReads	kmers	dup	cov	taxID	rank	taxName
4.7660	12155	5002	72665	1.45	1.342e-03	9785	species
0.3937	1004	9	3984	1.08	1.020e-04	9778	species
0.1380	352	25	1388	1.14	6.428e-04	9813	species
0.1313	335	3	1351	1.01	8.105e-07	9612	species
0.1686	430	430	1794	1.00	1.212e-06	9606	species
5.1680	13180	12862	178949	1.01	5.993e-03	1491	species

Filtering is critical for understanding content of ancient metagenomic samples!



KrakenUniq filters in aMeta: kmers = 1000, taxReads = 200

Pochon et al. *Genome Biology* (2023) 24:242
<https://doi.org/10.1186/s13059-023-03083-9>

Genome Biology

METHOD

Open Access

aMeta: an accurate and memory-efficient ancient metagenomic profiling workflow

Zoé Pochon^{1,2†}, Nora Bergfeldt^{1,3,4†}, Emrah Kırdök⁵, Mário Vicente^{1,2}, Thijessen Naidoo^{1,2,6,7}, Tom van der Valk^{1,4}, N. Ezgi Altınışık⁸, Maja Krzewińska^{1,2}, Love Dalén^{1,3}, Anders Götherström^{1,2†}, Claudio Mirabello^{9†}, Per Unneberg^{10†} and Nikolay Oskolkov^{11*†} 

[†]Zoé Pochon, Nora Bergfeldt, Anders Götherström, Claudio Mirabello, Per Unneberg, and Nikolay Oskolkov shared authorship.


*Correspondence: Nikolay.Oskolkov@bio.ltu.se

¹¹ Department of Biology, Science for Life Laboratory, National Bioinformatics Infrastructure Sweden, Lund University, Lund, Sweden
Full list of author information is available at the end of the article

Abstract

Analysis of microbial data from archaeological samples is a growing field with great potential for understanding ancient environments, lifestyles, and diseases. However, high error rates have been a challenge in ancient metagenomics, and the availability of computational frameworks that meet the demands of the field is limited. Here, we propose aMeta, an accurate metagenomic profiling workflow for ancient DNA designed to minimize the amount of false discoveries and computer memory requirements. Using simulated data, we benchmark aMeta against a current state-of-the-art workflow and demonstrate its superiority in microbial detection and authentication, as well as substantially lower usage of computer memory.

Keywords: Ancient metagenomics, Pathogen detection, Microbiome profiling, Ancient DNA



Where do they come from?

“For the discovery of pathogens in human patients ... a unique k-mer count threshold of 1000 eliminated many background identifications” (F. Breitwieser et al., 2018)

Comparing KrakenUniq filtering strategies



frontiers
in Microbiology

TYPE Original Research
PUBLISHED 18 May 2026
DOI 10.3389/fmicb.2026.1603339



OPEN ACCESS

EDITED BY
Ke Wang,
Fudan University, China

REVIEWED BY
Francisco Pascoal,
University of Porto, Portugal
Lucas Czech,
Kobenhavns Universitet Section for
GeoGenetics, Denmark

*CORRESPONDENCE
Nikolay Oskolkov
✉ nikolay.oskolkov@scilifelab.se

RECEIVED 31 March 2025
REVISED 20 April 2026
ACCEPTED 30 April 2026
PUBLISHED 18 May 2026

CITATION
Oskolkov N (2026) Refining filtering
criteria of Kraken family of tools for
accurate taxonomic profiling of ancient
metagenomic data.
Front. Microbiol. 17:1603339.
doi: 10.3389/fmicb.2026.1603339

COPYRIGHT
© 2026 Oskolkov. This is an


Refining filtering criteria of Kraken family of tools for accurate taxonomic profiling of ancient metagenomic data

Nikolay Oskolkov^{1,2*}


¹Department of Biology, National Bioinformatics Infrastructure Sweden, Science for Life Laboratory, Lund University, Lund, Sweden, ²Metabolic Research Group, Latvian Institute of Organic Synthesis, Riga, Latvia

Taxonomic profiling is a key component of ancient metagenomic analysis, however it is also susceptible to false-positive identifications. In particular, taxonomic classification tools from the Kraken family, such as Kraken2 and KrakenUniq, are highly sensitive to the choice of filtering options. To address this issue, various filtering approaches have been proposed. In this study, I conduct a comprehensive benchmarking of different filtering strategies for Kraken family of tools using simulated microbial and environmental ancient metagenomic data. I evaluate these approaches based on the balance between sensitivity and specificity of ground truth reconstruction (F1-score), and propose an optimal thresholding strategy tailored to specific sequencing depths in ancient metagenomic datasets.

Regular microbial dataset




Pathogenic microbial dataset



Three simulated datasets

env / sedaDNA dataset




Goal: reconstruct ground truth and find optimal Kraken filters

Optimal filters maximizing reconstruction accuracy


Jaccard similarity averaged across samples: regular microbes dataset

kmers_max = 1000, taxReads_max = 0, Jaccard_max = 0.785632805716957




Jaccard similarity averaged across samples: pathogen dataset


kmers_max = 1000, taxReads_max = 0, Jaccard_max = 0.564545454545454



Jaccard similarity averaged across samples: sedaDNA dataset

kmers_max = 1500, taxReads_max = 0, Jaccard_max = 0.736111111111111





Original KrakenUniq paper
(F. Breitwieser et al., Genome Biology 2018)

Metagenomics

Deduplicated files were merged by sequencing run and analyzed using Kraken2 [22]. We used a custom database containing all dusted bacterial, viral, archaeal, and protozoan reference sequences, as well as the UniVec_Core database. Kraken2 output files were filtered for 152 pathogenic taxa. Additionally, we computed the metagenomic profile for our sample using KrakenUniq [23]. Here, the custom database contained dusted complete genomes and chromosome-level assemblies of bacteria, viruses, archaea, and protozoa. Additionally, this database contained the human genome, the NCBI Viral Neighbor database, and the contaminant databases UniVec and EmVec. The KrakenUniq algorithm is better suited to identify low abundance hits and is able to identify unique *k*mers as well as ascertain the duplication rate and coverage of the *k*mer dictionaries of each taxon. To take advantage of this, we computed a heatmap using *plotly*, *pandas*, *matplotlib* [66], and *numpy* [67] based on an *E*-value calculated as follows: $\left(\frac{K}{R}\right) \times C$. Here, *K* is the *k*mer count, *R* is the read count, and *C* is the coverage of the taxon *k*mer dictionary. The *E*-value cutoff for further inspection was 0.001.

$$\text{E-value: } \left(\frac{\text{KmerCount}}{\text{ReadCount}}\right) \times \text{Coverage} = 0.001 - 0.1$$

Meriam Guellil et al., Genome Biology 2022

Maxime Borry

[Home](#) [Publications](#) [Courses](#) [Blog](#) [Contact me](#) [CV](#)

particularly problematic between scenario A and scenario B.

MODIFIED *DEXPE*-SCORE

Here I propose a slightly different score to mitigate some of issue. Instead of defining

$$E = \frac{K}{R} \times C$$

I propose to use







$$E = \frac{K}{R} \times \text{dexpe}(C)$$

with *dexpe* being the [double exponential function](#).







This way, more emphasis is given to the “coverage”, which particularly helps in low “coverage” scenarios, since the double exponential function grows “really fast” between 0 and 1.

Maxime Borry: <https://maximeborry.com/post/kraken-uniq>

Regular microbial dataset




Pathogenic microbial dataset




Environmental / sedaDNA dataset


KrakenUniq K filter
kmers_max = 1400, IoU_max = 0.7361111111111111




KrakenUniq C filter
C_max = 0.00065, IoU_max = 0.5208333333333333




KrakenUniq R filter
R_max = 360, IoU_max = 0.7361111111111111




KrakenUniq K / R filter
(K / R)_max = 13.5, IoU_max = 0.159756097560976





KrakenUniq (K / R) * C filter
((K / R) * C)_max = 0.001, IoU_max = 0.45



KrakenUniq (K / R) * dexp(C) filter
((K / R) * dexp(C))_max = 18, IoU_max = 0.1625





Take home messages:

- Probably the number of unique kmers alone is enough for optimal filtering
- Combinations of different filters do not seem to improve filtering quality

KrakenUniq vs. Kraken2 filters

Distinct minimizer count information

The [KrakenUniq](#) project extended Kraken 1 by, among other things, reporting an estimate of the number of distinct k-mers associated with each taxon in the input sequencing data. This allows users to better determine if Kraken's classifications are due to reads distributed throughout a reference genome, or due to only a small segment of a reference genome (and therefore likely false positive).

Thanks to the generosity of KrakenUniq's developer Florian Breilwieser in allowing parts of the KrakenUniq source code to be licensed under Kraken 2's MIT license, this distinct counting estimation is now available in Kraken 2. Development work by Martin Steinegger and Ben Langmead helped bring this functionality to Kraken 2.

At present, this functionality is an optional *experimental feature* -- meaning that we may later alter it in a way that is not backwards compatible with previous versions of the feature.

To use this functionality, simply run the `kraken2` script with the additional `--report-minimizer-data` flag along with `--report`, e.g.:

```
kraken2 --db $DBNAME --report k2_report.txt --report-minimizer-data \
--output k2_output.txt sequence_data.fq
```

This will put the standard Kraken 2 output (formatted as described in [Standard Kraken Output Format]) in `k2_output.txt` and the report information in `k2_report.txt`. Within the report file, two additional columns will be present, e.g.:

normal report format:

36.40	182	182	S2	211044	Influenza A virus (A/Puerto Rico/8/1934(H1N1))
-------	-----	-----	----	--------	--

modified report format:

36.40	182	182	1688	18	S2	211044	Influenza A virus (A/Puerto Rico/8/1934
-------	-----	-----	------	----	----	--------	---

In this modified report format, the two new columns are the fourth and fifth, respectively representing the number of minimizers found to be associated with a taxon in the read sequences (1688), and the estimate of the number of distinct minimizers associated with a taxon in the read sequence data (18). This would indicate that although 182 reads were classified as belonging to H1N1 influenza, only 18 distinct minimizers led to those 182 classifications.


The format with the `--report-minimizer-data` flag, then, is similar to that described in [Sample Report Output Format], but slightly different. The fields in this new format, from left-to-right, are:

1. Percentage of fragments covered by the clade rooted at this taxon
2. Number of fragments covered by the clade rooted at this taxon
3. Number of fragments assigned directly to this taxon
4. Number of minimizers in read data associated with this taxon (**new**)
5. An estimate of the number of distinct minimizers in read data associated with this taxon (**new**)
6. A rank code, indicating (U)nclassified, (R)oot, (D)omain, (K)ingdom, (P)hylum, (C)lass, (O)rder, (F)amily, (G)enus, or (S)pecies. Taxa that are not at any of these 10 ranks have a rank code that is formed by using the rank code of the closest ancestor rank with a number indicating the distance from that rank. E.g., "G2" is a rank code indicating a taxon is between genus and species and the grandparent taxon is at the genus rank.
7. NCBI taxonomic ID number
8. Indented scientific name

Is filtering by the number of distinct/unique minimizers in Kraken2 comparable with filtering by number of unique k-mers in KrakenUniq?



KrakenUniq / Kraken2 K filter

Optimal KrakenUniq kmers = 1100, optimal Kraken2 minimizers = 3300



KrakenUniq / Kraken2 K filter: pathogenic microbial dataset

Optimal KrakenUniq kmers = 900, optimal Kraken2 minimizers = 1700




Take home messages:


KrakenUniq and Kraken2 are comparable in accuracy provided optimal filters have been found for both

Kraken2 seems to have at least 2-3 times higher optimal number of distinct minimizers compared to optimal number of unique kmers for KrakenUniq

Optimal kmers vs. sequencing depth



“We found that the ideal thresholds increase by about 2000 unique k-mers per 1 million reads” (F. Breitwieser et al., 2018)



My experiment:

Seq Depth

100 000

500 000

700 000

Optimal n_kmers

200 + / - 295

1000 + / - 657

1400 + / - 945

$\text{optimal_n_unique_kmers} \sim 0.002 * \text{seq_depth}$

Approximately 200 unique kmers per 100 000 reads

My model prediction:

Seq Depth

1 000 000

5 000 000

10 000 000

Optimal n_kmers

2000

10000

20000

???


???

DATABASES!



DATABASES!

DATABASES!

DATABASES!





Why large databases are important for accurate taxonomic profiling



How can we use microbial abundance matrix from Kraken?

	sample1	sample2	sample3	sample4	sample5	sample6	sample7	sample8	sample9	sample10
<i>Ralstonia solanacearum</i>	3628	3751	619	1804	1384	1608	1375	0	1112	749
<i>Mycobacterium avium</i>	8236	8546	273	265	3221	4808	7750	6382	0	0
<i>Burkholderia pseudomallei</i>	7095	0	0	0	13082	0	4885	1456	0	7310
<i>Salmonella enterica</i>	4356	4471	4205	3588	0	13854	0	0	1959	2560
<i>Pseudomonas chlororaphis</i>	296	1024	0	977	374	677	276	0	0	294
<i>Neisseria meningitidis</i>	465	502	0	0	7341	0	0	3268	5643	0
<i>Yersinia pestis</i>	0	0	0	7174	957	0	11485	6461	0	11553



NIH Human Microbiome Project



Characterization of the microbiomes of healthy human subjects at five major body sites, using 16S and metagenomic shotgun sequencing.


Enter HMP1



Characterization of microbiome and human host from three cohorts of microbiome-associated conditions, using multiple 'omics technologies.

Enter IHMP



Enter HMP Integrated Portal  





Physalia
Courses

Understanding marker microbes for different environments

	Oral	Gut	Skin	Vagina
0	Neisseria	Blautia	Staphylococcus	Mobiluncus
1	Veillonella	Faecalibacterium	Peptoniphilus	Sphingopyxis
2	Actinomyces	Bacteroides	Citrobacter	Ureaplasma
3	Haemophilus	Dorea	Enhydrobacter	Caulobacter
4	Rothia	Akkermansia	Finegoldia	Gardnerella
5	Leptotrichia	Clostridium	Propionibacterium	Chlamydia
6	Cardiobacterium	Ruminococcus	Acinetobacter	Asticcacaulis
7	Capnocytophaga	Subdoligranulum	Massilia	Mycobacterium
8	Oribacterium	Oxalobacter	Hymenobacter	Herbaspirillum
9	Alloprevotella	Oscillibacter	Corynebacterium	Lactobacillus
10	Gemella	Eubacterium	Bacillus	Achromobacter
11	Fusobacterium	Bilophila	Micrococcus	Atopobium



Microbial source tracking



Duitama González et al. *Microbiome* (2023) 11:243
<https://doi.org/10.1186/s40168-023-01670-3>

Microbiome

METHODOLOGY

Open Access

decOM: similarity-based microbial source tracking of ancient oral samples using k-mer-based methods

Camila Duitama González^{1*}, Riccardo Vicedomini^{1,2}, Téio Lemane², Nicolas Rascovan³, Hugues Richard⁴ and Rayan Chikhi¹


Abstract

Background The analysis of ancient oral metagenomes from archaeological human and animal samples is largely confounded by contaminant DNA sequences from modern and environmental sources. Existing methods for Microbial Source Tracking (MST) estimate the proportions of environmental sources, but do not perform well on ancient metagenomes. We developed a novel method called decOM for Microbial Source Tracking and classification of ancient and modern metagenomic samples using k-mer matrices.

Results We analysed a collection of 360 ancient oral, modern oral, sediment/soil and skin metagenomes, using stratified five-fold cross-validation. decOM estimates the contributions of these source environments in ancient oral metagenomic samples with high accuracy, outperforming two state-of-the-art methods for source tracking, FEAST and mSourceTracker.

Conclusions decOM is a high-accuracy microbial source tracking method, suitable for ancient oral metagenomic data sets. The decOM method is generic and could also be adapted for MST of other ancient and modern types of metagenomes. We anticipate that decOM will be a valuable tool for MST of ancient metagenomic studies.

Keywords Ancient metagenomics, Microbial source tracking, k-mer matrix, Paleogenomics





*Knut och Alice
Wallenbergs
Stiftelse*



LUNDS
UNIVERSITET